

Response to the Associate Editor and Reviewers

Manuscript: *FlowMOP: An Automated Flow Cytometry Time, Debris, and Doublet Removal Tool*

The Associate Editor and reviewers raised several overlapping concerns. To avoid repeating the same response while ensuring that every point is addressed, closely related comments are reproduced verbatim and answered together below. A complete comment-coverage index is provided at the end of this response.

We thank the Associate Editor and reviewers for their considered feedback and believe that the requested amendments have significantly improved the manuscript. We hope that the revised manuscript now addresses the Associate Editor's and reviewers' concerns.

Associate Editor's General Assessment

Novel tools to expedite flow cytometry data analysis are always welcome, and I commend the authors for their effort in introducing these new methods. However, as currently submitted, the manuscript requires major revision. The reviewers have identified numerous areas requiring improvement, and I encourage the authors to address their comments carefully and comprehensively. In addition to the reviewers' critiques, I offer several observations of my own that I believe will strengthen the manuscript's suitability for publication.

Response: The revisions are organized around the algorithmic motivation, comparative mechanism, computational benchmarking, interpretation of expert evaluations, methodological clarity, and the limits of the current approach. Detailed responses follow.

Combined Comment 1 — Algorithmic Motivation, Novelty, Related Tools, and Validation Framing

Raised by: Associate Editor; Reviewer 2, Comments 1, 6, and 8.

Associate Editor — Algorithmic Motivation and Theoretical Justification

The authors acknowledge the existence of other algorithms that perform substantially similar analyses to the components of FlowMop, yet they do not adequately motivate what deficiencies in the existing algorithms their new approach is designed to address. A compelling framing would follow a structure such as: we demonstrated that PeacoQC and FlowCut fail under realistic conditions X, Y, and Z; upon analysis, we determined this is attributable to specific aspects of how those algorithms operate; accordingly, we designed our algorithm to overcome these limitations. No such explanation is currently provided, nor is there a clear justification for why the particular algorithmic choices were made.

Reviewer 2, Comment 1

In the abstract, the authors claim “FlowMOP is the first automated approach capable of debris cleaning”. However, there are methods like GateNet, 10.1002/cyto.a.20531, and UNITO that can do automatic gating and probably can be used to gate out debris. There is even a review for this purpose (10.3389/fimmu.2015.00380).

Reviewer 2, Comment 6

The introduction did not introduce available related tools.

Reviewer 2, Comment 8

The authors claimed, “Conversely, this paper seeks to compare performance not only through traditional human-expert defined standards, but also the development of bespoke data generated explicitly for pre-processing validation.” In fact, both simulation data and read data are commonly used for new method validation.

Response

We now frame FlowMOP as a training-free preprocessing workflow that combines automated time, debris, and doublet removal. We revised the Introduction to distinguish automated population-gating tools, time-quality-control tools, and FlowMOP's intended role as an integrated preprocessing workflow. We removed the claim that FlowMOP is the first automated approach capable of debris cleaning. The revised Introduction discusses FlowCut and PeacoQC as time-dependent quality-control approaches and GateNet and UNITO as broader automated population-gating approaches.

We have also revised the description of the three FlowMOP modules. Time gating is based on positive-population fluorescence summaries across acquisition order, debris gating derives an FSC-A threshold from cross-parameter positive-population structure, and doublet gating uses FSC-A/FSC-H and SSC-A/SSC-H ratio structure. These algorithmic choices are now distinguished from computational implementation choices.

In response to Reviewer 2's observation that simulated and real data are both commonly used in method validation, we now describe the synthetic samples specifically as event-labelled technical controls developed to provide objective ground truth for flow-cytometry preprocessing tasks.

Changes made

- We removed the claim that FlowMOP is the first automated approach capable of debris cleaning and added the following clarification:

“Several tools automate cytometry population identification, including the model-based flowClust method [13], the neural-network method GateNet [14], and the bivariate-segmentation framework UNITO [15]; automated gating approaches are reviewed elsewhere [16]. FlowMOP is a Python-based, training-free preprocessing workflow that combines automated time-gating, debris removal, and doublet exclusion in a single headless tool. FlowCut and PeacoQC provide the most direct comparison for its time-gating component.”

- We reframed the synthetic-data contribution as follows:

“The synthetic datasets provide event-level labels for estimating retained-target purity and removed-non-target purity in preprocessing tasks where real ground truth is otherwise difficult to define.”

- We added a concise description of the three algorithmic components to the Figure 1 legend:

“A) FlowMOP selects valid parameters using positive-peak detection, generates

time-binned fluorescence summaries, applies two smoothing resolutions, and performs robust outlier rejection across acquisition order. B) FlowMOP applies the same valid-parameter check, derives a candidate FSC-A threshold from each eligible parameter’s positive events, and uses the median candidate as the final FSC-A gate. C) FlowMOP identifies doublets from inflection points in FSC-A/FSC-H and SSC-A/SSC-H ratio histograms, with a fixed-ratio fallback.”

Location: Introduction; Results, “Algorithmic design,” “Time Gating,” “Debris Gating,” and “Doublet Gating”; Figure 1 legend.

Combined Comment 2 — Computational Benchmarking and Within-Sample Parallelisability

Raised by: Associate Editor; Reviewer 1, Comment 1; Reviewer 2, Comment 7.

Reviewer 1, Comment 1 — Technical Infrastructure and Performance Claims

The adoption of a Dask-based Python framework is a highly commendable architectural choice. Flow cytometry analysis has historically under-utilized distributed compute resources, and providing a modern, scalable alternative to R-based packages is a significant contribution to the “Big Data” era of spectral cytometry. However, the manuscript currently lacks empirical data to support the claims of improved run-time or memory efficiency, while other contemporary algorithms explicitly highlight their scalability for large datasets.

Recommendation: To make a compelling case for FlowMOP’s adoption, the authors could include benchmarking figures (e.g., execution time and peak memory usage) comparing FlowMOP against PeacoQC and FlowCut across datasets of increasing file count and/or size ($\sim 10^3$ to $\sim 10^7$ events or some [Event x Parameter] matrix varying each). Another point of comparison may be in the distributed compute scenario vs. more traditional local compute resources.

Example Performance Benchmarking Table:

Dataset Size (Events)	Metric	FlowMOP (Python/Dask)	PeacoQC (R)	FlowCut (R)
10^4	Execution			
	Time (s)			
	Peak RAM (MB)			
10^5	Execution			
	Time (s)			
	Peak RAM (MB)			
...etc.				

Associate Editor — Algorithmic vs. Implementation Advantages

Reviewer 1 raises several points regarding the use of the Dask framework, and I would like to extend that line of inquiry. Specifically, I would ask the authors to clarify whether there is something inherent to the FlowMop algorithm that makes

it particularly well-suited to parallelization via Dask — as distinct from the implementation itself. This distinction matters considerably. Given current development tools, a competent programmer with AI tooling could likely port the existing FlowCut or PeacoQC code — neither of which is especially lengthy — to a Dask-based parallel implementation in a matter of hours. If that is the case, it raises the question of whether a substantial portion of FlowMop's reported advantage derives from the implementation rather than from any intrinsic algorithmic innovation, which would significantly diminish the novelty of the contribution. I would encourage the authors to focus their comparative analysis on algorithmic advantages, particularly in cases where existing algorithms present genuine structural barriers to parallelization.

Reviewer 2, Comment 7

The authors should introduce “Dask” before using the “Dask-based framework” in the introduction. Please keep in mind, most readers probably have no experience using Python at all. I use Python a lot myself. Still, I am not familiar with this “Dask” thing. Also, the Dask-based design is about calculation efficiency, instead of accuracy. Readers will be more interested in which algorithm is used for automatic gating, instead of how to speed up calculation. The authors did not really introduce the algorithm for gating in the introduction. But they use a whole paragraph on the Dask-based design.

Response

We evaluated computational performance using matched 36-channel FCS inputs containing 10,000, 100,000, 300,000, 1,000,000, and 2,000,000 events. FlowMOP, PeacoQC, and FlowCut were run on the same clone-based inputs. Optional plotting, reporting, and output generation were disabled where supported, and FlowMOP's debris and doublet modules were disabled so that the shared time-gating task was timed fairly. Each condition included one warm-up and three measured repeats. Runtime and peak resident memory are reported as mean $\pm$ SD in Table 1. At 2,000,000 events, FlowMOP was approximately 2.5-fold faster than PeacoQC and 3.5-fold faster than FlowCut while using approximately 32% and 45% of their respective peak RAM.

The benchmark used local non-distributed execution. Testing active Dask execution showed that its overhead did not improve this workload, so Dask was inactive for the reported benchmark and removed from the manuscript's novelty framing.

Among the three evaluated decision structures, only FlowMOP is inherently suited to independent within-sample channel parallelisation. Once shared acquisition bins have been defined, each eligible channel independently establishes its fluorescence reference, produces a fixed-shape time-bin summary, and applies smoothing and MAD filtering. Cross-channel coordination occurs only when the resulting bin flags are combined in the final parameter vote. This creates coarse, regular channel-level tasks with a single final reduction.

PeacoQC must instead reconcile peak configurations across bin-channel combinations before its MAD, Isolation Tree, and consecutive-bin decisions. FlowCut combines segment- and channel-level statistics through adjacent-segment and file-wide comparisons, contiguous-region decisions, and an optional flagged-file rerun. Both methods can process separate files in parallel, and some internal suboperations could be parallelised, but their complete within-sample decision pipelines contain intermediate dependencies that prevent the same independent channel-wise reduction without changing their decision rules. Reimplementation in Python or another scheduling frame-

work would not remove those structural dependencies.

FlowMOP's within-sample parallelisability is therefore an inherent property of its decision structure rather than its programming language. The reported runtime benchmark used local non-distributed execution, so it measures the performance of the tested implementations rather than an empirical distributed-computing speedup.

Changes made

- We removed Dask from the manuscript's novelty and accuracy framing and specified the benchmark configuration as follows:

“For fair timing of the shared time-gating task, FlowMOP was run using local non-distributed execution, with debris and doublet removal disabled and annotated output FCS writing disabled. PeacoQC and FlowCut were run with optional plotting, reporting, and output generation disabled where supported.”

- We state FlowMOP's contribution and identify the most direct time-gating comparators:

“FlowMOP is a Python-based, training-free preprocessing workflow that combines automated time-gating, debris removal, and doublet exclusion in a single headless tool. FlowCut and PeacoQC provide the most direct comparison for its time-gating component.”

- We retained a concise operational description in the Methods:

“FlowMOP accepts `.csv`, `.fcs`, and Parquet files. It reduces event-level measurements into time-bin and channel summaries, applies smoothing and robust outlier detection, and combines the resulting flags through parameter voting.”

- We placed the architectural comparison and its relationship to the computational benchmark in the Discussion:

“These measurements were obtained using local non-distributed execution and therefore demonstrate the performance of the tested implementations, not a distributed-computing speedup.”

“Among the three evaluated decision structures, only FlowMOP exposes an inherently channel-parallel within-sample computation. After shared acquisition bins are defined, each eligible fluorescence channel independently establishes its reference, produces a fixed-shape time-bin summary, and applies smoothing and MAD filtering; cross-channel coordination occurs only in the final parameter vote. PeacoQC instead performs data-dependent peak identification and reconciliation across bin-channel combinations, while FlowCut applies adjacent-segment, contiguous-region, file-wide, and conditional rerun decisions [5,9]. Their suboperations and separate files can be parallelised, but their complete within-sample decision pipelines cannot be expressed as the same independent channel-wise reduction without changing their decision rules. Porting either implementation to another language or scheduling framework would therefore not remove these structural dependencies. This structural distinction, together with FlowMOP's earlier data reduction, fewer full-data passes, and smaller intermediate state, is consistent with its lower runtime and memory use in our benchmarks, although

the benchmark does not independently establish causation.”

- We added matched runtime and peak-memory testing across five event counts:

“Computational scalability was evaluated using clone-based real-FCS scaling. A representative FCS file was subsampled/replicated to matched event counts of 10,000, 100,000, 300,000, 1,000,000, and 2,000,000 events while preserving the original 36-channel structure. FlowMOP, PeacoQC, and FlowCut were run on the same generated inputs for each size.”

- We reported the principal computational result as follows:

“The computational benchmark demonstrates that FlowMOP provides speed gains with lower peak RAM usage at larger event counts. This is most evident from 300,000 events onward, where FlowMOP had both the fastest mean runtime and lowest peak memory use. At 2,000,000 events, FlowMOP was approximately 2.5-fold faster than PeacoQC and 3.5-fold faster than FlowCut, while using approximately 32% of PeacoQC’s peak RAM and 45% of FlowCut’s peak RAM.”

Location: Methods, “Algorithmic design” and “Computational scalability”; Results, “Computational scalability”; Table 1; Discussion, “Synthetic Sample Time Gating.”

Combined Comment 3 — Gating Mechanism, Smoothing, and Parameter Fairness

Raised by: Associate Editor; Reviewer 1, Comment 5; Reviewer 2, Comment 20.

Associate Editor — Algorithmic Motivation and Theoretical Justification

More troublingly, the authors themselves appear uncertain about why their algorithm performs better in certain cases. In the Discussion, they write that a performance difference “may be attributed to FlowMop’s ‘smoothing’ implementation.” This speculation is unsatisfying — can the authors not test this directly by modifying the smoothing implementation and evaluating whether the results change accordingly? While presenting cherry-picked examples in which one algorithm outperforms another has illustrative value, a rigorous justification of the theoretical underpinnings of the algorithmic design is essential. Without it, readers cannot reasonably assess which sample types and experimental conditions are expected to favor FlowMop over existing alternatives.

Reviewer 1, Comment 5 — Parametric Sensitivity Analysis

To strengthen the claim of “superiority,” it is essential to demonstrate that the results are not merely a function of suboptimal default settings for the competing algorithms.

Recommendation: A supplemental figure showing a sensitivity analysis—where the parameters of PeacoQC and FlowCut are tuned across a range—would provide a more “apples-to-apples” comparison and confirm that FlowMOP’s performance gains are algorithmic rather than parametric.

Reviewer 2, Comment 20

The authors say, “Finally, FlowMOP also novelly employs the use of smoothing at multiple resolutions to ensure multiple types of aberrations are detected”. What are the “multiple types of aberrations”? please clarify.

Response

We now seek to address the mechanisms underlying FlowMOP's superior performance under the primary fixed comparison settings in the tested Segment, Bimix, and Trimix scenarios. Our hypotheses arose from the different signals used by the methods and from the pattern observed in the original benchmark. FlowMOP uses acquisition order to define time bins but bases its removal decision on fluorescence summaries of globally defined positive populations; Time itself is excluded from the quality variables. FlowCut, by contrast, can respond to local acquisition-density structure. This distinction was particularly relevant because the Segment files retained stronger source-linked Time-density structure and showed the clearest FlowCut performance loss, whereas Time was approximately normalized in the constructed Bimix and Trimix files. We therefore hypothesised that FlowCut's relative loss arose, at least in part, from sensitivity to local acquisition rate rather than to the source-linked fluorescence defect itself.

To test this hypothesis, we conducted a matched Time-only mechanism analysis. This involved changing only the Time channel while preserving fluorescence, scatter, source labels, and event order across 30 source-labelled Bimix, Trimix, and Segment files. We found that FlowMOP and PeacoQC were both unchanged under source-linked and random Time warping, whereas FlowCut's removal behaviour changed when local Time density was altered. The clearest loss occurred in Segment inputs under source-linked Time warping.

PeacoQC's documentation identifies acquisition-bin size as a trade-off between the accuracy of within-bin density estimation, the number of bins available for evaluating signal stability, and the number of events affected when a bin is removed [5]. The Bimix and Trimix benchmarks contain short, interspersed source-defined intervals. FlowMOP's stronger performance under the fixed benchmark settings is therefore consistent with its temporal summarisation being better matched to these brief intervals, whereas PeacoQC's automatically selected bin resolution reflects a general-purpose trade-off.

Consequently, the earlier smoothing attribution has been removed. The matched Time-only benchmark directly supports the distinction between FlowMOP's fluorescence-population-summary decision and FlowCut's response to acquisition-density changes when fluorescence is unchanged. PeacoQC's documented bin-size trade-off provides the relevant context for its fixed-setting comparison with FlowMOP.

We also tested smoothing directly by holding all other FlowMOP settings fixed across all 92 large-cut source-labelled synthetic time-gating files (Table S4). Relative to no smoothing, the current dual-resolution default (0.01, 0.05) increased mean retained-target purity by 3.1%, while mean removed-non-target purity decreased by 3.6%. Across the tested dual settings, 0.02, 0.09 produced the highest retained-target purity. Thus, dual-resolution smoothing improved retained-target purity, with a corresponding removed-non-target purity trade-off.

During revision, we also corrected the terminology used for these two source-composition measures. The quantity previously called “sensitivity” is retained target-source events divided by all retained events, and the quantity previously called “specificity” is removed non-target-source events divided by all removed events. Because these are retained-set and removed-set purity

measures rather than conventional sensitivity and specificity, we now call them retained-target purity and removed-non-target purity throughout the manuscript, tables, and figure labels. This relabelling does not alter the values, comparisons, or p-values.

The two smoothing resolutions are now described as complementary spline fits applied to the same time-bin fluorescence-summary series before MAD filtering. A bin can be detected through either smoothing pass; the manuscript no longer relies on the imprecise phrase “multiple types of aberrations.”

Changes made

- We removed the earlier smoothing attribution and added the following matched comparison:

“To test the effect of flow-rate disturbances aligned with or independent of source-linked fluorescence changes, we altered only the Time channel while leaving fluorescence, scatter, source labels, and event order unchanged (Fig. S4). Flow-MOP and PeacoQC were unchanged under both source-linked and random Time warping. In contrast, FlowCut's retained-target and removed-non-target purity shifted after Time-only perturbation.”

- We clarified the function of the two smoothing resolutions as follows:

“Two spline smoothing values, one small and one larger (current default 0.01, 0.05), are applied to the returned time-bin series before median absolute deviation (MAD) filtering.”

“The two smoothing resolutions target both shorter and more sustained deviations, while parameter voting limits removal driven by isolated noisy channels in higher-dimensional panels.”

- We added a direct smoothing ablation across all 92 largecut synthetic files as Supplementary Table S4. The manuscript reports the tested dual-resolution grid, the no-smoothing control, the former primary-comparison setting (0.1, 0.9), and the current software default (0.01, 0.05).
- We corrected the metric terminology and added the following definitions:

“Retained-target purity was defined as retained target-source events divided by all retained events. Removed-non-target purity was defined as removed non-target-source events divided by all removed events.”

- We added the complete mechanism-benchmark methods and results as Figure S4.

Location: Methods, “Time-only acquisition-rate mechanism benchmark” and “MAD-smoothing ablation and default selection”; Results, “Time Gating” and “Synthetic Time Gating Benchmark”; Figure S4; Supplementary Table S4; Discussion, “Synthetic Sample Time Gating.”

Combined Comment 4 — Parameter Voting and the Ten-Parameter Threshold

Raised by: Reviewer 2, Comments 18 and 21.

Reviewer 2, Comment 18

The authors say "when >10 parameters are present, bins are rejected if they are flagged by two or more parameters". Why "10" is used as threshold here?

Reviewer 2, Comment 21

The authors say, "parameter voting to ensure that higher dimensional sample sets are not overly degraded." It would be nice to show how bad the overly degradation will be is parameter voting is not applied. It is a little a bit concern that events with aberrant signal in one fluorescence are not detected, even in a big panel.

Response

The revised manuscript identifies the ten-parameter cutoff as an empirical operational threshold rather than a mathematically derived boundary.

The rationale for requiring two channel flags in panels with more than ten eligible parameters is conservative. Biological datasets have a non-zero probability of isolated channel-specific outliers, and the opportunity for such outliers increases with panel dimensionality. Requiring corroboration from a second parameter reduces the risk that a single noisy channel removes otherwise valid biological events. The corresponding trade-off is that an abnormality confined to one fluorescence channel may be retained in a high-dimensional panel. We did not add a separate voting-ablation analysis and have avoided presenting the rule as universally optimal.

Changes made

- We added the following rationale and limitation:

"For panels with more than 10 parameters, FlowMOP requires two or more parameters to flag a bin before rejection. This empirical safeguard reduces false-positive removal caused by isolated noisy channels in high-dimensional panels, although an aberration confined to one channel may consequently be retained (Figure 1A)."
- We did not add a voting-ablation experiment and do not present the threshold as universally optimal.

Location: Results, "Time Gating."

Combined Comment 5 — Expert Rankings, Terminology, Visualisation, and Bayesian Modelling

Raised by: Associate Editor; Reviewer 1, Comment 3; Reviewer 2, Comment 16.

Associate Editor — Expert Preference Rankings and Statistical Interpretation

My other major concern pertains to the substantial space — both in the main manuscript and supplemental materials — devoted to preferential rankings by four experts across multiple sample types and algorithms, analyzed using a Bayesian statistical framework. Setting aside the question of whether all experts can reasonably be considered equally expert across all sample types, and the limited statistical power afforded by four data points, I find the interpretation of these rankings fundamentally ambiguous.

As expected, experts generally preferred human-gated data over automated tools — but this finding is difficult to interpret in isolation. Forcing experts to rank outputs necessarily produces an ordering, but that ordering conflates two very different scenarios: one in which all automated results are unacceptable and the other in which all automated results are perfectly adequate, with human gating simply preferred at the margin. Similarly, even if FlowMop was ranked above the other automated tools, the reader has no basis for judging whether this difference reflects a meaningful improvement in analytical utility or merely a marginal statistical preference with no practical consequence.

I would suggest that the authors replace or supplement the ranking analysis with a simpler, more interpretable quality scale — for example: 1 = excellent gating; 2 = adequate and suitable for analysis; 3 = inadequate for analysis but not terrible; 4 = unacceptable. This approach would convey absolute quality rather than relative preference and would be more informative to readers.

Reviewer 1, Comment 3 — Refinement of Human Subjective Rankings

While the subjective rankings provide an interesting window into expert perspectives on automated gating, the presentation and terminology in this section require significant refinement:

Nomenclature: In computer science, the term “benchmarking” typically implies an objective, repeatable measure of performance in specific scenarios. In the context of subjective human preference, “Comparison” or “Evaluation” is a more accurate/appropriate term. This sentence, “FlowMOP’s outputs in these samples were benchmarked against a set of human experts”, should more precisely say that FlowMOP’s outputs were “compared against expert gating”.

Methodological Robustness: Utilizing a single expert per tissue type limits the statistical power of the ranking methodology. It prevents the reader from seeing the intra-tissue consensus—how consistently multiple experts on the same tissue agree on the rank-order of methods used on a specific sample. What it currently highlights is how a single tissue expert has a unique perspective of the correct gating for a given tissue and inter-operator gating variability; not algorithmic gating superiority. Other improvements could be adding different processing conditions (e.g. fixation) or more variation in tissue used (e.g. include human PBMC or other species’ tissues).

Data Visualization: The ranking scale (7 = Best, 1 = Worst) is possibly counter-intuitive and should be reversed to follow standard ordinal ranking conventions. The row axis title in the grid figure should be updated to “Gate Provided By” to clearly disambiguate between the creator of the gate (e.g., “Expert 1”) and the individual performing the ranking (“Expert”).

The scatter plot in Panel B adds limited value; replacing it with an “Average Score” column at the end of the grid would be cleaner. This statement is oddly contorted to be true, but is misconstruing the data: “However, it is of note that in 4/5 and 7/9 datasets in the debris and doublets tasks respectively, FlowMOP was not the least preferred, indicating superiority to at least one human benchmarker (Fig. 6A, 7A)”. Another interpretation is that FlowMOP had the lowest average score in Fig. 7A/B and had the fewest tissue cases where experts deemed it as the best result or even better-than-average traditional gating.

Reviewer 2, Comment 16

There is only a “Bayesian Modelling” section under the “Non-synthetic samples” section in the Method. The Bayesian Modelling section described the method to rank performance of human experts, FlowMOP, FlowCut and PeacoQC. It is not about how “Non-synthetic samples” are generated at all. Also, the purpose of the Bayesian Modelling is well described in the beginning of this subsection, which makes it difficult to follow.

Response

Forced rankings measure relative preference and cannot determine whether a gate is absolutely acceptable for analysis. We therefore replaced “benchmarking” with “expert comparison” or “expert evaluation,” removed the contorted “not least preferred” interpretation, and now state directly where FlowMOP ranked below the human gates. The ranking results are presented as exploratory evidence of expert preference rather than objective ground-truth evidence of gating quality or algorithmic superiority.

The design used a single relevant expert to rank each tissue type and therefore cannot estimate within-tissue consensus across multiple raters. We did not retrospectively collect a new absolute-quality scale because that would require re-running the expert assessment under a newly defined protocol. The revised manuscript states this limitation and interprets the rankings alongside the synthetic technical-control results that contain event-level ground truth.

The visual summaries now include an Average Score column within each ranking grid rather than a separate panel, and the row axis is labelled “Gate Provided By.” The displayed scale now follows the standard convention that rank 1 is best. This display correction did not change the underlying ranking order or Bayesian analysis.

The Bayesian model is retained because the observed outcome is an ordinal ranking. The Plackett–Luce model retains the complete ordering and represents uncertainty in relative preference without treating differences between adjacent ranks as equal-interval continuous measurements. We have clarified that it estimates relative preference only. The subsection is presented as the statistical analysis of the expert-ranking data rather than as a method for generating non-synthetic samples.

Changes made

- We defined the scope of the expert-ranking analysis as follows:

“The expert-ranking analysis was used to summarize relative preferences among gates generated by FlowMOP, comparator algorithms, and human operators; it was not treated as an absolute measure of gating adequacy.”

- We revised the Results terminology and interpretation as follows:

“FlowMOP’s outputs in these samples were compared with expert-provided gates using a forced-ranking preference task. These rankings measure relative expert preference and should not be interpreted as an absolute measure of gating adequacy.”

- We added the following limitations:

“One relevant expert ranked each tissue type, only four experts contributed gates, and the ordering does not distinguish a marginal preference from a judgement that a gate is unsuitable for analysis. Gate style could also reveal which outputs were algorithmic. We therefore interpret the ranking results as exploratory relative preferences, not as absolute quality scores or proof of algorithmic superiority or inferiority.”

- We retained the Plackett–Luce analysis and clarified that it models ordinal preference rankings rather than absolute gating quality.
- We regenerated Figures 5–7 as editable SVGs, changed the row-axis title to “Gate Provided By,” displayed rank 1 as best, and replaced the separate summary panel with an Average Score column in each grid.
- We revised the Figure 5 caption as follows:

“Figure 5. Expert preference rankings for time gates provided by four human experts, FlowMOP (black border), FlowCut, and PeacoQC across nine datasets. Rows indicate the gate provider, and the final column shows the mean rank across datasets. Rank 1 indicates greatest preference; therefore, a lower average score indicates greater preference. Abbreviations: DRG, dorsal root ganglion; CNS, central nervous system.”

- We clarified the scoring direction in the Results as follows:

“Rank 1 indicates the greatest preference.”

Location: Methods, “Statistical Analysis of Expert Preference Rankings”; Results, “Expert Preference Evaluation” and Figures 5–7; Discussion, “Biological Datasets: Expert Preference Evaluation.”

Combined Comment 6 — Downstream Biological Impact and the Cost of Errors

Raised by: Associate Editor; Reviewer 1, Comment 4.

Associate Editor — Expert Preference Rankings and Statistical Interpretation

Alternatively — or in addition — the authors could evaluate the downstream biological impact of each gating strategy. By substituting FlowMop, FlowCut, and PeacoQC gates for the time and debris gating steps while retaining the human gating strategy for all other steps, one could directly assess whether differences in automated gating performance have any meaningful effect on the biological conclusions drawn from the data. Ultimately, the relevant question for end users is not whether an algorithm removes a few percent more or fewer debris events, but whether doing so materially affects downstream analysis.

Reviewer 1, Comment 4 — Discussion of Error “Costs” and Sensitivity Trade-offs

The manuscript notes that PeacoQC exhibits higher sensitivity but lower specificity compared to FlowMOP. This trade-off is central to clinical discovery and warrants a deeper discussion.

Discussion Point: The authors should explicitly address the biological “cost” of these errors. For instance, is it more detrimental to the study's integrity to leave in a small transient artifact or to accidentally remove a rare, clinically significant biological population?

Response

Neither higher removal nor closer agreement with a single human gate is inherently preferable. Under-cleaning can retain acquisition artifacts that create, inflate, or obscure apparent biological populations. Conversely, over-cleaning can remove rare or transient biological events. The revised Discussion therefore interprets retained-target purity and removed-non-target purity as complementary composition measures with competing error costs rather than treating one direction of error as universally preferable. These terms replace the earlier nonstandard use of “sensitivity” and “specificity”: retained-target purity is the proportion of retained events from the target source, whereas removed-non-target purity is the proportion of removed events from non-target sources.

PLACEHOLDER

We added an initial downstream biological-concordance analysis using three tumour and three non-tumour human liver samples. We compared the complete FlowMOP time, debris, and doublet intersection with the matched final manual gate and used an operational Zombie UV-A threshold as an orthogonal readout. To distinguish biological selectivity from general permissiveness, the primary endpoint was the Zombie-high/Zombie-low removal-rate ratio. The matched FlowMOP/manual comparison was tested using an exact two-sided Wilcoxon signed-rank test on the log-transformed paired ratios.

PLACEHOLDER

FlowMOP showed greater selectivity in four of six samples, with a median paired FlowMOP/manual fold-change of 1.68 (range 0.82-2.70), but the paired difference was not statistically significant ($p = 0.156$). We interpret this as no detected systematic difference in this small sample set, not as proof of equivalence. A secondary decomposition showed that time gating was approximately Zombie-neutral, debris removal was heterogeneous, and the doublet mask was the principal source of Zombie-high enrichment. A further biological validation against a prespecified population or functional endpoint remains a PLACEHOLDER and will be added before submission.

Changes made

- We added the following discussion of error costs:

“The biological cost of preprocessing errors is difficult to measure directly, and neither under-cleaning nor over-cleaning is preferable. Under-cleaning may allow acquisition-time artifacts with abnormal staining patterns to confound downstream results, including by creating, inflating, or obscuring apparent rare populations. Conversely, over-cleaning could remove rare or transient biological populations.”

- We added a paired tumour-liver biological-concordance analysis and report the complete per-sample results in Supplementary Tables S3A-B.

- We used a permissiveness-adjusted primary endpoint and a matched statistical test:

“FlowMOP had greater permissiveness-adjusted Zombie-high removal selectivity than manual gating in four of six samples. The median paired FlowMOP/manual selectivity fold-change was 1.68 (range 0.82-2.70), but the difference was not statistically significant (exact two-sided Wilcoxon signed-rank test on log ratios, $p = 0.156$).”

- A second downstream biological validation remains a PLACEHOLDER in this response and is not presented as completed manuscript content.

Location: Methods, “Tumour-liver acquisition-instability and biological-concordance analysis”; Results, “Tumour-liver acquisition instability and downstream biological concordance”; Discussion, “Biological Datasets: Expert Preference Evaluation”; Supplementary Tables S3A-B.

Combined Comment 7 — Dataset Scale, Complexity, and Tumor Validation

Raised by: Associate Editor; Reviewer 2, general assessment and Comment 10.

Associate Editor — Smaller Point

In the Introduction, the authors state that sophisticated analysis software has contributed to datasets of increasing size and complexity. It is not clear to me how analysis software makes debris or time gating more complex, particularly when pre-processing is predominantly scatter-based. Do the authors perhaps mean that derived parameters from imaging cytometry increase pre-processing complexity? Similarly, the claim that “datasets exceed human capabilities” presumably refers to the sheer number of files rather than the complexity of individual datasets — this should be stated explicitly. It is also not apparent from the manuscript how the tools described here depend on data complexity per se.

Reviewer 2 — General Assessment

The authors developed An Automated Flow Cytometry Time, Debris, and Doublet Removal Tool. This can be a useful tool for cytometry. However, the dataset used in the paper are relatively “easy” cases. I recommend the authors to challenge the FlowMOP with more difficult but common samples like digested tumor sample. Performance against human expert is needed. Also, the manuscript is not well structured. Careful edition by experienced researcher is needed to increase the readability of the manuscript.

Reviewer 2, Comment 10

The study did not validate the FlowMOP on tumor datasets, which is a common sample type for flow cytometry but difficult to perform debris removal. I highly suggest the authors to include a tumor dataset, even if the performance is not optimal, the readers can know whether the method can be applied in their sample or not.

Response

The revised Introduction distinguishes study scale from the intrinsic complexity of an individual preprocessing gate. It now refers explicitly to increasing numbers of files, events, and measured parameters and explains that this scale makes repeated manual preprocessing time-consuming

and often impractical. It no longer implies that analysis software itself necessarily makes a debris or time gate more complex.

We added an analysis of three tumour and three non-tumour human liver samples as a difficult validation setting. For the tumour files, we first evaluated acquisition-time instability by comparing the intervals detected by FlowMOP with manual gating, FlowCut, and PeacoQC and by examining fluorescence-only divergence. FlowMOP captured the principal high-divergence acquisition intervals identified by the other approaches and also identified additional intervals with positive-population fluorescence shifts; residual comparator-only intervals generally showed weaker fluorescence divergence.

We then evaluated the complete time, debris, and doublet workflow using Zombie UV-A as an orthogonal biological readout. FlowMOP and manual gating were not detectably different in the primary paired selectivity comparison (exact two-sided $p = 0.156$). This result does not establish formal parity, but it provides no evidence of systematic inferiority in this six-sample analysis. We retain an explicit limitation that larger and more necrotic tumour cohorts are needed because these samples lack event-level technical ground truth and cannot represent the full diversity of tumour-digest artifacts.

Changes made

- We clarified the distinction between study scale and individual-gate complexity as follows:

“Modern flow-cytometry studies can contain increasing numbers of files, events, and measured parameters [2]. The large scale of these data makes repeated manual preprocessing time-consuming and often impractical, potentially amplifying variability between operators. Reproducible automated preprocessing is therefore valuable for large studies.”

- We added a tumour-liver validation comprising three tumour and three non-tumour samples.
- We added fluorescence-only acquisition-instability analysis for the tumour files and a paired manual-versus-FlowMOP downstream Zombie analysis for all six liver samples.
- We revised the remaining limitation as follows:

“The present liver analysis provides an initial test in tumour tissue, but further tumour digests with severe necrosis, large aggregates, and heterogeneous scatter profiles remain important validation contexts because they may challenge a conservative FSC-A-centered debris strategy.”

Location: Methods and Results, “Tumour-liver acquisition instability and biological concordance”; Discussion, “Synthetic Sample Debris and Doublet Gating” and “Biological Datasets: Expert Preference Evaluation”; Supplementary Tables S3A-B.

Combined Comment 8 — Acquisition Settings and Newer Scatter Configurations

Raised by: Reviewer 2, Comments 13 and 28.

Reviewer 2, Comment 13

The voltage set of the cytometer significantly impact the resolution of debris. If the voltage is relatively low with bad resolution, can FlowMOP perform the automatic task well? Also, please describe the voltage setting in the data generation section in the Method.

Reviewer 2, Comment 28

The Xenith flow cytometry form Thermo Fisher provides 405 FSC/SSC, 488 FSC/SSC, and polar 488 FSC/SSC. This provides more flexible approach to gate target population. How can the FlowMOP adapt to this kind of flow data? It would be nice to discuss this in the Discussion.

Response

The acquisition voltage/gain settings remain a PLACEHOLDER.

The present validation does not cover all configuration choices introduced by instruments that offer multiple scatter measurements. FlowMOP currently expects the user to identify the scatter channels used by the workflow; it has not been validated systematically across 405-nm, 488-nm, and polar scatter configurations. Adaptation to these instruments therefore remains a future validation and channel-selection task.

Changes made

- The acquisition voltage/gain information remains a PLACEHOLDER.
- We added the following scope statement:

“Instruments providing multiple scatter measurements, including 405-nm, 488-nm, and polar 488-nm FSC/SSC configurations, were not evaluated here. Future versions could assess multiple scatter-channel pairs and select or combine the pair with the clearest debris/doublet separation.”

Location: Discussion, “Other remarks”; acquisition voltage/gain information: PLACEHOLDER.

Combined Comment 9 — Temporal Artifacts and Synthetic Time-Sample Design

Raised by: Associate Editor; Reviewer 2, Comments 3, 4, and 14.

Associate Editor — Smaller Point

Regarding the bimix and trimix data used to simulate realistic fluorescence fluctuations: I am uncertain how these mixtures reproduce the kinds of flow rate variations that occur during actual experiments. If I understand correctly, the simulated fluctuations were subtle enough that human reviewers could not reliably detect them. I would ask the authors to clarify whether this represents a clinically or experimentally common problem — in my experience, flow rate deviations typically manifest as clear departures at the beginning or end of an acquisition, with mid-acquisition instabilities being relatively uncommon.

Reviewer 2, Comment 3

The authors should explain the concept of “temporal artifacts” in the abstract, given this is the core issue to be addressed in the study.

Reviewer 2, Comment 4

There are many terms in the abstract “(concatenated and mixed time perturbations, high-debris mixtures, and CFSE/CTV co-labeled doublets)”. I am not sure whether it is a good idea to put them here without explanation, considering unexplained term might add confusion to the readers. Anyway, the authors should introduce these terms in the introduction, which is missing.

Reviewer 2, Comment 14

The “Generation of Synthetic Time Samples” section is difficult to understand. This section is important for readers to understand the validation results. I recommend the authors use a illustration to show these 3 synthetic manners. Also, please explain why these three manners are designed in this way? So that the readers can follow the text.

Response

The revised text explains the practical relevance of the Bimix and Trimix designs. It defines a “microblockage” operationally as a short, self-resolving mid-acquisition disturbance that produces a localized fluorescence shift; the term does not imply that a physical obstruction was directly observed. Existing work supports the substance of this phenomenon even though it uses broader terminology. flowAI documented flow-rate surges interspersed throughout acquisitions and their association with signal-intensity variation [10]. PeacoQC describes temporary acquisition shifts caused by slow sample uptake or a clog and notes that such problems can be difficult to detect manually [5], while FlowCut describes manual identification and removal of transient acquisition problems as time-consuming and subjective [9].

Because the affected intervals can be short and their events interspersed with otherwise plausible events, they are difficult to recognize reliably by eye and impractical to exclude using a series of manual gates. We use “microblockage” for this subtle, self-resolving acquisition failure mode that a time-quality-control method should be able to detect. Estimating its prevalence across experiments would require a separate study.

The Bimix and Trimix samples model the observable fluorescence consequence in this operational definition rather than reproducing or proving a physical obstruction. This distinction is important. The synthetic samples can appear acceptable by eye, but their retained source labels reveal short intervals containing events from an intentionally perturbed fluorescence source. Under the benchmark definition, those events should be excluded even when visual inspection alone would not identify a defensible manual gate. The apparent visual normality of these files is therefore the reason an event-labelled benchmark is needed, not evidence that the modeled problem is unimportant.

The biological validation further supports the experimental relevance of this design: the tumour-liver files contained mid-acquisition intervals with coordinated multichannel fluorescence shifts, including subtle intervals detected by FlowMOP in addition to the principal disturbances shared with manual gating or comparator methods. PLACEHOLDER: the additional biological validation will be inserted here once complete. Together, the synthetic and biological analyses

motivate a benchmark that includes both readily visible sustained artifacts and visually subtle transient microblockages.

We also defined temporal artifacts, simplified the terminology, and added Figure S1. We clarified that the primary synthetic samples model fluorescence changes across acquisition order without flow-rate disturbances, because rate changes without corresponding fluorescence changes should not prompt event exclusion. Flow-rate effects were tested separately, either aligned with source-linked fluorescence changes or independently of them. FlowMOP was unchanged in both conditions, whereas FlowCut's removal behavior shifted (Figure S4).

Changes made

- In the Abstract, we added the following definition:

“Temporal artifacts are acquisition-time-dependent deviations in event quality or fluorescence signal caused by blockages, bubbles, flow instability, or instrument instability.”

- In the Methods, we added the following clarification:

“Here, we use ‘microblockage’ operationally to denote a short, self-resolving mid-acquisition disturbance that produces a localized fluorescence shift; the term does not imply that a physical obstruction was directly observed. Segmented samples model sustained changes, whereas Bimix and Trimix model the observable fluorescence consequence in this operational definition by introducing short source-defined fluorescence shifts during acquisition. They do not recreate or establish the physical mechanism itself. The Bimix and Trimix files can appear acceptable on visual inspection because the altered intervals are short and interspersed with otherwise plausible events; however, the retained source labels identify the intentionally perturbed events that should be excluded under the benchmark definition.”

“Flow-rate disturbances without corresponding fluorescence changes should not prompt event exclusion; these samples therefore model fluorescence changes across acquisition order without introducing flow-rate disturbances. Flow-rate effects were tested separately by altering Time either in alignment with source-linked fluorescence changes or independently of them.”

- In the Results, we added the following comparison:

“To test the effect of flow-rate disturbances aligned with or independent of source-linked fluorescence changes, we altered only the Time channel while leaving fluorescence, scatter, source labels, and event order unchanged (Fig. S4). FlowMOP was unchanged under both source-linked and random Time warping. In contrast, FlowCut's retained-target and removed-non-target purity shifted after Time-only perturbation.”

- In the Discussion, we now explain why the visually subtle synthetic cases are experimentally relevant and require an event-labelled benchmark:

“Transient acquisition disturbances are not confined to the beginning or end of a run: flow-rate surges interspersed throughout acquisition and associated signal-

intensity variation have been documented [10]. PeacoQC notes that temporary acquisition problems can be difficult to detect manually [5], and FlowCut describes manual identification and removal of transient acquisition problems as time-consuming and subjective [9]. We use ‘microblockage’ operationally for a short, self-resolving instance of this broader phenomenon that produces a localized fluorescence shift, without asserting that a physical obstruction was directly observed.”

“Although these synthetic samples can appear acceptable on visual inspection, the source labels show that the short altered intervals contain events from an intentionally perturbed fluorescence source and therefore should be excluded under the benchmark definition. Visual subtlety is thus a central feature of the benchmark: it demonstrates why apparent normality by eye is not sufficient ground truth.”

- We linked this rationale to the biological validation:

“The tumour-liver validation further supports the experimental relevance of this design by demonstrating mid-acquisition intervals with coordinated multichannel fluorescence shifts.”

- We added Figure S1 with the following caption:

“Figure S1: Construction of Segment, Bimix, and Trimix synthetic time samples. No flow-rate disturbance was introduced. Source labels were retained for scoring only and excluded from algorithm input.”

Location: Abstract; Introduction; Methods, “Generation of Synthetic Time Samples” and “Time-only acquisition-rate mechanism benchmark”; Results, “Synthetic Time Gating Benchmark” and “Tumour-liver acquisition instability and downstream biological concordance”; Discussion, “Synthetic Sample Time Gating”; Figures S1 and S4.

Combined Comment 10 — Abstract and Introduction Structure

Raised by: Reviewer 2, Comments 2 and 9.

Reviewer 2, Comment 2

The structure of abstract is kind of odd. The first paragraph should focus on background instead of providing a conclusion.

Reviewer 2, Comment 9

The introduction is missing an end paragraph for summary.

Response

The Abstract has been restructured to begin with the preprocessing problem and its biological and computational context before introducing FlowMOP, the validation design, and the principal findings. We also added a concluding Introduction paragraph that summarizes FlowMOP's intended role and previews the time-gating comparison, debris and doublet validation, expert evaluation, and computational benchmark.

Changes made

- We reordered the Abstract so that it begins with the study context:

“Flow cytometry now generates high-parameter datasets whose scale and variability challenge manual preprocessing, leading to subjectivity and poor reproducibility. Temporal artifacts are acquisition-time-dependent deviations in event quality or fluorescence signal caused by blockages, bubbles, flow instability, or instrument instability. This study introduces FlowMOP, a Python-native framework that automates three major preprocessing steps—time-gating, debris removal, and doublet exclusion.”

- We added the following study summary to the end of the Introduction:

“Here, we introduce FlowMOP, an automated preprocessing tool for time-gating, debris removal, and doublet exclusion. We compare time-gating performance with PeacoQC and FlowCut, evaluate debris and doublet removal using synthetic ground-truth datasets and expert comparison, and benchmark computational scalability.”

Location: Abstract; final paragraph of the Introduction.

Combined Comment 11 — Debris Methodology and Validation

Raised by: Reviewer 1, Comment 2; Reviewer 2, Comments 15, 22, and 24.

Reviewer 1, Comment 2 — Debris Gating Methodology

The inclusion of an automated debris-gating method is a welcome addition, particularly as this is often omitted in other QC tools. However, the current reliance on FSC-A as the primary determinant for debris may explain the lower rankings in expert consensus (Figure 6A).

Critique: While the rationale is that debris is generally “low-size,” this does not account for larger morphological debris, such as large-cell clumps (e.g. tumor) or non-biological aggregates, which can occupy higher FSC/SSC ranges. Metadata-based margin removal techniques may augment this debris filtering.

Recommendation: The authors should provide a more robust rationale for leaning exclusively on FSC-A or, ideally, discuss how the integration of SSC-A or Pulse Width might improve debris discrimination beyond ultra-low-size events.

Reviewer 2, Comment 15

The authors say “For assessment of FlowMOP gating performance, ‘high debris’ and ‘low debris’ samples were synthetically combined.” How are they combined? Please remember, well-described Method guarantees the transparency of the study.

Reviewer 2, Comment 22

The FlowMOP applies an FSC-A based threshold to exclude debris. I have concern about this FSC-A based idea. Commonly, we use FSC vs. SSC scatter to gate debris, which is clear in samples like PBMC. However, the resolution of a simple FSC-A

histogram will not be as good as that in a FSC vs. SSC scatter, because debris and cells will overlap in FSC-A axis. In the Figure 3, the authors show representative results which is based on simple samples, which is possible to use the FSC-A based threshold method to gate debris. However, in difficult cases, which is not very rare, this FSC-A based threshold method may perform very bad.

Reviewer 2, Comment 24

I would like to see the debris remove rate calculated against the human expert.

Response

The revised Methods describes the debris-removal procedure more precisely. FlowMOP ultimately applies a one-dimensional FSC-A threshold, but that threshold is not obtained from a single unconditioned FSC-A histogram. The algorithm identifies eligible fluorescence parameters, examines the FSC-A structure of each parameter's positive population, derives a candidate FSC-A threshold from each eligible channel, and applies the median candidate threshold to the sample. Thus, information across eligible fluorescence channels informs the final FSC-A gate, while the applied decision remains a conservative FSC-A threshold.

The current module is not a universal debris classifier. Debris and intact cells can overlap on FSC-A, and large clumps, tumour-associated material, non-biological aggregates, or other high-scatter contaminants may not be removed reliably. We specifically considered whether SSC-A should be incorporated. SSC-A can be informative for large or internally complex debris, but large and internally complex events can also represent desirable populations that should be retained. Their interpretation varies with tissue, staining panel, cytometer configuration, acquisition settings, and the biological populations of interest. Unlike the comparatively transferable relationship between very low FSC-A and small debris, there is no broadly safe high-SSC-A rule for excluding large events.

In the present technical-control datasets, the labelled small-debris and desirable source populations overlapped substantially along SSC-A, so a fixed SSC-A threshold offered limited additional separation for the small-event removal targeted by FlowMOP. We therefore did not incorporate SSC-A into the current debris decision. Pulse-width measurements may similarly help identify some aggregates or clumps, but their availability and interpretation are instrument- and acquisition-dependent. Incorporating SSC-A or pulse width responsibly would require a wide configurable parameter range or sample-specific multivariate decision rules validated for the intended dataset, panel, instrument, and cell populations. The revised Discussion therefore defines the current scope as conservative low-FSC debris removal and presents SSC-A, pulse-width measurements, margin-event metadata, and sample-specific multivariate models as future extensions rather than universal default filters.

The synthetic preparation is now described explicitly. Approximately equal event numbers were sampled from the high-debris and low-debris sources and concatenated into matched mixtures while retaining source labels. These source labels provide the objective basis for the reported post-cleaning proportions.

The requested human comparison is already represented by the percentage-based analysis in Figure 3. FlowMOP's enrichment of the labelled low-debris component is compared directly with four expert gates, with mean percentages, variability, and paired statistical comparisons. Because the input event count differs between samples, percentage-based retention and enrichment are more directly comparable than raw removed-event counts.

Changes made

- We described the synthetic debris mixture as follows:

“For assessment of FlowMOP debris-gating performance, high-debris and low-debris samples were combined by sampling approximately equal event numbers from each source and concatenating them into matched synthetic mixtures while retaining source labels for ground-truth quantification.”

- We clarified the scope and derivation of the FSC-A gate as follows:

“To debris gate, FlowMOP applies a conservative FSC-A-based threshold intended primarily to remove low-FSC debris. FSC-A was selected because low-forward-scatter material is a comparatively universal debris signal across sample types, whereas SSC-A patterns are more tissue- and instrument-dependent.”

“SSC-A can provide useful complementary information about internal complexity and may help identify some large debris or aggregates. However, large, high-SSC-A events can also be desirable intact populations, and their interpretation depends on the tissue, staining panel, cytometer configuration, and acquisition settings. Unlike the relatively transferable rule that very small, low-FSC events are likely to contain debris, there is no comparably universal SSC-A direction or threshold for removing large events. In the present technical-control datasets, the labelled small-debris and desirable source populations also showed substantial overlap along SSC-A, so a fixed SSC-A rule provided limited additional discrimination for the small-event removal targeted by FlowMOP. SSC-A was therefore not incorporated into the current debris decision. A future SSC-A implementation would require configurable or sample-specific multivariate decision rules validated across the intended tissues, panels, and instruments.”

“The median FSC-A threshold across all parameters is taken as the final FSC-A gate to be applied to the sample (Figure 1B).”

- We expanded the Discussion of SSC-A and pulse-width integration:

“SSC-A may improve recognition of some large or internally complex debris, and pulse-width measurements may help distinguish some aggregates or clumps. However, large and internally complex events can also be biologically desirable, while pulse-width availability and behaviour vary between instruments and acquisition configurations. Consequently, a fixed high-SSC-A or pulse-width exclusion rule could remove legitimate populations, and an appropriate decision boundary would need to be tailored to the dataset, panel, instrument, and intended biological populations.”

- We reported the existing expert comparison as follows:

“FlowMOP did not differ significantly from any human evaluator except Expert 4, for whom FlowMOP removed more labelled high-debris events (Bonferroni-adjusted paired t-test, $p = 0.04$) (Fig. 3C).”

- We identified potential extensions for complex debris phenotypes without claiming that they are part of the current implementation:

“Future extensions could incorporate SSC-A, pulse-width measurements, margin-event metadata, or sample-specific multivariate models when these features are available and appropriately validated.”

- We added the following limitation:

“The present liver analysis provides an initial test in tumour tissue, but further tumour digests with severe necrosis, large aggregates, and heterogeneous scatter profiles remain important validation contexts because they may challenge a conservative FSC-A-centered debris strategy.”

Location: Methods, “Synthetic Debris Sample Preparation and Generation”; Results, “Debris Gating” and “Synthetic Debris Gating Benchmark”; Discussion, “Synthetic Sample Debris and Doublet Gating”; Figure 3.

Combined Comment 12 — Aggregates, Doublets, Saturation, and Validation

Raised by: Reviewer 2, Comments 5, 25, and 26.

Reviewer 2, Comment 5

Preprocessing sometimes include aggregates removal. Is it possible for FlowMOP to handle this task?

Reviewer 2, Comment 25

The FlowMOP creates a histogram of the FSC-A/FSC-H ratio to gate doublets. I also have a concern about this method. One unmentioned hypothesis of this algorithm is that all cells fall within the measurement range of FSC. However, it is very common that some cells exceed the FCS limit and collapsed in the edge of a FSC-A vs FSC-H scatter plot. For example, big myeloids in PBMC, if the target population is lymphocytes. Such population is expected to lead to weird histogram of the FSC-A/FSC-H ratio. This can be difficult to be addressed by the FlowMOP and are not mentioned in the study.

Reviewer 2, Comment 26

The authors use CTV-CFSE double positive cells to represent doublets. However, there are probably CTV- CTV doublets and CFSE - CFSE doublets, which are missed in the performance evaluation. I would like to see the doublet remove rate calculated against the human expert.

Response

FlowMOP's doublet module targets events whose FSC-A/FSC-H and, where available, SSC-A/SSC-H pulse ratios distinguish them from singlets. It can remove some higher-order coincident events when those events produce separable ratio structure, as observed for triplet-like events in the technical controls, but it cannot identify every aggregate or clump from these ratios alone.

Informative pulse ratios require the relevant scatter measurements to remain within the instrument's measurement range. If cells are saturated or collapsed at an FSC boundary, the lost pulse-shape information cannot be reconstructed downstream. The revised Results now state

this assumption and identify saturated or edge-collapsed scatter and incorrectly configured acquisition parameters as conditions requiring acquisition review, manual intervention, or alternative pulse-shape features. We did not add a new saturation-handling algorithm.

The reviewer is also correct that CTV-CFSE double-positive events represent only heterologous labelled doublets. CTV-CTV and CFSE-CFSE doublets cannot be distinguished from their respective single-labelled populations using the dye labels alone. The technical-control design therefore gives a sensitive measure of retained known heterologous doublets but cannot establish the complete false-positive and false-negative rate for all doublet classes.

Figure 4 already compares the percentage of CTV-CFSE double-positive events removed by FlowMOP with the corresponding percentages removed by human experts. We now describe this endpoint and its same-label limitation explicitly rather than implying complete doublet ground truth.

Changes made

- We added the following measurement-range limitation:

“FlowMOP's doublet module assumes that FSC-A/FSC-H and SSC-A/SSC-H ratios remain informative and that acquisition voltages and scatter parameters have been set appropriately. If relevant scatter channels are saturated, edge-collapsed, or incorrectly configured at acquisition, the lost pulse-shape information cannot be recovered by FlowMOP or by other downstream preprocessing algorithms; such samples require acquisition review, manual intervention, or alternative pulse-shape features where available.”

- We clarified the limits of the CTV/CFSE validation as follows:

“Events positive for both CFSE and CTV provide an observable class of heterologous labelled doublets, subject to rare dye-transfer events; same-label CTV-CTV and CFSE-CFSE doublets are not identifiable from these labels alone.”

- We reported the comparison with human experts as follows:

“No statistically significant difference was detected between FlowMOP and any expert for this endpoint (Fig. 4B, paired t-test; unadjusted $p > 0.05$).”

- We retained the observation that FlowMOP removed a triplet-like population in the technical-control samples and restricted that finding to those samples.

Location: Results, “Doublet Gating” and “Synthetic Doublet Gating Benchmark”; Discussion, “Synthetic Sample Debris and Doublet Gating” and “Other remarks”; Figure 4.

Combined Comment 13 — CTV/CFSE Preparation Protocol

Raised by: Reviewer 2, Comment 27.

Reviewer 2, Comment 27

The protocol about how the CTV-CFSE samples are prepared is not clear. Are there wash steps in between? This is very important. It would be nice if the authors can provide detailed protocol as supplementary files.

Response

PLACEHOLDER

Changes made

PLACEHOLDER

Location: PLACEHOLDER

Combined Comment 14 — Missing and Duplicated Methods Information

Raised by: Reviewer 2, Comments 11 and 12.

Reviewer 2, Comment 11

There is duplicated line in the Method “Samples were acquired on a Cytex Northern Lights 3-laser (V/B/R) spectral flow cytometer.”

Reviewer 2, Comment 12

In the “Preparation of Human PBMC Samples for Synthetic Time Benchmarking Samples” section, the authors mentioned there are many antibodies in the staining cocktail. However, they did not list which antibodies.

Response

The duplicated cytometer sentence has been removed.

PBMC antibody identities and fluorochromes: PLACEHOLDER

Changes made

- We removed the duplicated sentence, “Samples were acquired on a Cytex Northern Lights 3-laser (V/B/R) spectral flow cytometer.”
- The complete PBMC antibody panel remains a PLACEHOLDER.

Location: Methods, “Preparation of Human PBMC Samples for Synthetic Time Benchmarking Samples”; complete PBMC antibody panel: PLACEHOLDER.

Combined Comment 15 — Precleaning Event Removal and Threshold

Raised by: Reviewer 2, Comment 17.

Reviewer 2, Comment 17

In the “Precleaning” section of Results, how many events are being removed in these used datasets? This information can be provided as supplementary table/figure to justify the set of threshold (5%).

Response

The current implementation uses a 1% threshold, not 5%; the inconsistent earlier value has been corrected. FlowMOP checks the number of events at the maximum FSC-A value and removes those maximum-valued events only when they exceed 1% of the sample.

The number and percentage of events removed by this specific precleaning step are the relevant quantities. A dataset-level summary of those values is not yet complete and remains to be added before submission; percentage-based retention results from later gating stages are not presented as a substitute for this precleaning summary.

The 1% cutoff is a pragmatic safeguard rather than a uniquely ground-truth-derived boundary. It is low enough to identify files with a non-trivial accumulation of events at the acquisition maximum while avoiding removal triggered by isolated maximum-valued events. The revised Results state this operational rationale.

Changes made

- We corrected and defined the threshold as follows:

“FlowMOP first checks the input file for events at the limit of detection, defined here as events at the maximum FSC-A value for that sample. If the number of events at this maximum exceeds a threshold (default 1%), FlowMOP removes these maximum-valued events. Otherwise, it retains all values.”

- We added the following operational rationale:

“The 1% cutoff is a pragmatic safeguard intended to identify non-trivial accumulation at the acquisition maximum without responding to isolated maximum-valued events.”

- The dataset-level precleaning removal counts and percentages remain to be added before submission.

Location: Results, “Precleaning.”

Combined Comment 16 — Figure 1 Axes and Annotations

Raised by: Reviewer 2, Comments 19 and 23.

Reviewer 2, Comment 19

The Figure 1A missing many axis titles. Also, there are tow subplots after time-binned median fluorescence. What are these two subplots are not clearly annotated.

Reviewer 2, Comment 23

The Figure 1B also misses many axis titles. The Figure 1C misses y-axis title.

Response

Figure 1 remains a conceptual schematic, and its axes are not treated as quantitative measurements. The revised legend explains the two smoothing resolutions in Figure 1A, identifies

the debris and doublet-ratio quantities in Figures 1B and 1C, and states that the schematic fluorescence-intensity and signal-strength axes use arbitrary units.

Changes made

- We revised the Figure 1 legend as follows:

“Figure 1. Conceptual schematics depicting FlowMOP’s time-gating (A), debris-gating (B), and doublet-gating (C) methods; plotted fluorescence intensity and signal-strength axes are schematic and use arbitrary units. A) FlowMOP selects valid parameters using positive-peak detection, generates time-binned fluorescence summaries, applies two smoothing resolutions, and performs robust outlier rejection across acquisition order. B) FlowMOP applies the same valid-parameter check, derives a candidate FSC-A threshold from each eligible parameter’s positive events, and uses the median candidate as the final FSC-A gate. C) FlowMOP identifies doublets from inflection points in FSC-A/FSC-H and SSC-A/SSC-H ratio histograms, with a fixed-ratio fallback.”

Location: Figure 1 legend.

Associate Editor's Closing Request

Finally, I ask that the authors also provide thorough and complete responses to all comments raised by the other reviewers.

Response: Every comment from Reviewers 1 and 2 is reproduced in this response, and the coverage index below identifies the consolidated response containing each comment. Completed revisions are described and quoted where available; author-supplied methodological details and the ongoing biological-validation analysis remain clearly marked as placeholders.

Location: Not applicable.

Comment Coverage Index

Associate Editor

Associate Editor point	Response location
General assessment	Associate Editor's General Assessment
Algorithmic motivation and theoretical justification — existing alternatives and design rationale	Combined Comment 1
Algorithmic motivation and theoretical justification — smoothing mechanism	Combined Comment 3
Algorithmic versus implementation advantages	Combined Comment 2
Expert preference rankings and statistical interpretation	Combined Comments 5 and 6
Smaller point — dataset scale and complexity	Combined Comment 7
Smaller point — Bimix/Trimix and experimental relevance	Combined Comment 9
Request for complete responses	Associate Editor's Closing Request

Reviewer 1

Reviewer 1 comment	Response location
Comment 1 — Technical infrastructure and performance claims	Combined Comment 2
Comment 2 — Debris gating methodology	Combined Comment 11
Comment 3 — Refinement of human subjective rankings	Combined Comment 5
Comment 4 — Error costs and sensitivity trade-offs	Combined Comment 6
Comment 5 — Parametric sensitivity analysis	Combined Comment 3

Reviewer 2

Reviewer 2 comment	Response location
General assessment	Combined Comment 7
Comment 1	Combined Comment 1
Comment 2	Combined Comment 10
Comment 3	Combined Comment 9
Comment 4	Combined Comment 9
Comment 5	Combined Comment 12
Comment 6	Combined Comment 1
Comment 7	Combined Comment 2
Comment 8	Combined Comment 1
Comment 9	Combined Comment 10
Comment 10	Combined Comment 7
Comment 11	Combined Comment 14
Comment 12	Combined Comment 14
Comment 13	Combined Comment 8
Comment 14	Combined Comment 9
Comment 15	Combined Comment 11
Comment 16	Combined Comment 5
Comment 17	Combined Comment 15
Comment 18	Combined Comment 4
Comment 19	Combined Comment 16
Comment 20	Combined Comment 3
Comment 21	Combined Comment 4
Comment 22	Combined Comment 11
Comment 23	Combined Comment 16
Comment 24	Combined Comment 11
Comment 25	Combined Comment 12
Comment 26	Combined Comment 12
Comment 27	Combined Comment 13
Comment 28	Combined Comment 8

Closing Statement

The revisions and explicit limitations provide a clearer account of the manuscript's evidentiary basis and practical interpretation.